

CONTENTS

DECLARATION	i
DEDICATION	ii
ACKNOWLEDGMENT	iv
ABBREVIATIONS	v
1 INTRODUCTION	1
1.1 Background	1
1.2 Motivation and significance	1
1.3 Problem statement	2
1.4 Aim and objectives	3
1.4.1 General objective	3
1.4.2 Specific objectives	3
1.5 Research questions	4
2 LITERATURE REVIEW	5
2.1 Introduction	5
2.2 Hospital overcrowding and triage delay	5
2.3 SATS in LMIC and Ghanaian practice	6
2.4 TEWS and vital-sign triage	6
2.5 Clinical pathways as operational roadmaps	7
2.6 Attention, BERT, and BioBERT	7
2.7 NLP for emergency and hospital triage	8
2.8 Conversational agents and digital triage tools	9
2.9 Summary: what prior work supplies and what it leaves open	9
2.10 Research gap	10
3 METHODOLOGY	12
3.1 Methodological overview	12
3.1.1 Chief complaints	13
3.1.2 Training data provenance	13

3.1.3	BERT and bidirectional context	14
3.1.4	BioBERT for biomedical text	14
3.1.5	The [CLS] summary token	14
3.1.6	Pipeline summary	14
3.1.7	Safety and nurse authority	16
3.2	SATS colour chart	17
3.3	TEWS: Triage Early Warning Score	17
3.3.1	Definition	18
3.3.2	Piecewise scoring functions	18
3.3.3	Colour map	19
3.3.4	Relationship to the text Softmax path	20
3.4	Neural networks, Transformers, and BERT	21
3.5	Introducing BERT	22
3.6	Introducing BioBERT	23
3.7	Tokenisation	23
3.7.1	Input X	23
3.7.2	The tokenisation map τ	23
3.7.3	Why the maximum length is $M = 128$	24
3.8	The [CLS] and [SEP] tokens	25
3.9	Embedding lookup	26
3.9.1	Why three embedding components	27
3.9.2	Why the hidden dimension is $D = 768$	28
3.10	Input matrix $H^{(0)}$	29
3.11	Encoder layers	30
3.12	Self-attention	32
3.13	Classification softmax	34
3.13.1	Mutually exclusive acuity categories	34
3.13.2	Softmax as the probability translator	34
3.13.3	Categorical distribution ($K = 5$)	35
3.13.4	Bernoulli as the $K = 2$ building block	35
3.14	Model training pipeline	37
3.14.1	Stage 1: BioBERT pre-training (masked language modelling)	37
3.14.2	Stage 2: Encoder transfer from a triage checkpoint	38
3.14.3	Stage 3: Fine-tuning on nurse-labeled chief complaints	38
3.15	Mapping acuity to colour: f_{SATS}	41
3.16	Clinical pathways $P(C_{\text{NLP}})$	41
3.17	Medical relevance gate	45
3.18	Worked numerical walkthrough	46
3.19	End-to-end summary	47

3.20	Implementation note	49
4	RESULTS	50
4.1	BioBERT model analysis	50
4.2	Fine-tuned model analysis	53
4.2.1	Evaluation protocol	53
4.2.2	Interpreting high holdout accuracy	54
4.2.3	SMOTE for class imbalance	55
4.2.4	Baseline fine-tuned model	56
4.2.5	SMOTE-augmented variant	56
4.2.6	Per-class and headline comparison	57
4.3	Inference latency	60
4.4	Pathway routing at KNUST University Hospital	60
4.4.1	Nurse triage vs chatbot response	60
4.5	Anticipated impact on healthcare delivery	62
4.5.1	Functional chatbot and workflow improvement	62
4.5.2	Reduced initial triage wait times	63
4.5.3	No overlooked critical cases	64
4.5.4	Patient satisfaction and appropriate categorisation	64
4.5.5	Staff workload and public-health insights	64
4.5.6	Success criteria	64
4.6	Discussion of findings	65
5	CONCLUSION	66
5.1	Limitations and future work	66
	REFERENCES	69
	APPENDIX	70
5.2	Software repository	70
5.3	Prototype /predict response fields	70
5.4	Stage 3 fine-tuning checklist	70
5.5	Glossary of principal symbols	71